

# Report

October 19, 2022

## 1 University of Central Florida

### 1.1 Department of Computer Science

### 1.2 CDA 5106: Spring 2020

### 1.3 Machine Problem 1: Cache Design, Memory Hierarchy Design by

### 1.4 « Sahar Sheikholeslami »

#### 1.4.1 Honor Pledge: “I have neither given nor received unauthorized aid on this test or assignment.”

1.4.2 Student’s electronic signature: \_\_\_\_\_ Sahar Sheikholeslami \_\_\_\_\_

(sign by typing your name)

```
[1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import math
```

## 2 1) GRAPH #1 (total number of simulations: 55)

2.0.1 For this experiment:

2.0.2 o L1 cache: SIZE is varied, ASSOC is varied, BLOCKSIZE = 32.

2.0.3 o L2 cache: None.

2.0.4 o Replacement policy: LRU

2.0.5 o Inclusion property: non-inclusive

2.0.6 Plot L1 miss rate on the y-axis versus  $\log_2(\text{SIZE})$  on the x-axis, for eleven different cache sizes:

2.0.7 SIZE = 1KB, 2KB, ..., 1MB, in powers-of-two. (That is,  $\log_2(\text{SIZE}) = 10, 11, \dots, 20$ .) The graph should contain five separate curves (i.e., lines connecting points), one for each of the following associativities: direct-mapped, 2-way set-associative, 4-way set-associative, 8-way setassociative, and fully- associative. All points for direct-mapped caches should be connected with a line, all points for 2- way set-associative caches should be connected with a line, etc.

```
[2]: BockSize = 32
size1 = []
log1 = []
fullyAssoc1=[]
for i in range (11):
    sizei = (2**i)*1024
    size1.append(sizei)
    logi= math.log(sizei, 2)
    log1.append(logi)
    fullyassoci=sizei/BockSize
    fullyAssoc1.append(fullyassoci)

    print("size of cache ", i+1, " is " , sizei, " and the log is: ", logi)

print(size1)
print(log1)
print( fullyAssoc1)
```

```
size of cache 1 is 1024 and the log is: 10.0
size of cache 2 is 2048 and the log is: 11.0
size of cache 3 is 4096 and the log is: 12.0
size of cache 4 is 8192 and the log is: 13.0
size of cache 5 is 16384 and the log is: 14.0
size of cache 6 is 32768 and the log is: 15.0
size of cache 7 is 65536 and the log is: 16.0
size of cache 8 is 131072 and the log is: 17.0
```

```

size of cache 9 is 262144 and the log is: 18.0
size of cache 10 is 524288 and the log is: 19.0
size of cache 11 is 1048576 and the log is: 20.0
[1024, 2048, 4096, 8192, 16384, 32768, 65536, 131072, 262144, 524288, 1048576]
[10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0, 17.0, 18.0, 19.0, 20.0]
[32.0, 64.0, 128.0, 256.0, 512.0, 1024.0, 2048.0, 4096.0, 8192.0, 16384.0,
32768.0]

```

## 2.0.8 L1 Miss Arrays for the 11 sizes and different Assoc

```

[3]: L1MissRates_Assoc1 =[0.193460,0.147740,0.100170,0.067000,0.046090,0.037680,0.
↪032920, 0.032330, 0.025840, 0.025840, 0.025840]
L1MissRates_Assoc2 =[0.156030,0.107140, 0.075280, 0.047340, 0.033840, 0.028810,↪
↪0.027130, 0.025900, 0.025840, 0.025820, 0.025820]
L1MissRates_Assoc4 =[0.142700, 0.096220, 0.059920, 0.042470, 0.028320, 0.
↪026400, 0.025950, 0.025820, 0.025820, 0.025820, 0.025820]
L1MissRates_Assoc8 =[0.136270,0.090690 , 0.053650, 0.039540, 0.027740, 0.
↪026250,0.025890,0.025820, 0.025820, 0.025820, 0.025820]
L1MissRates_FullyAssoc = [0.136960, 0.088600, 0.049540, 0.039120, 0.026340, 0.
↪026240, 0.025830,0.025820, 0.025820, 0.025820, 0.025820]

```

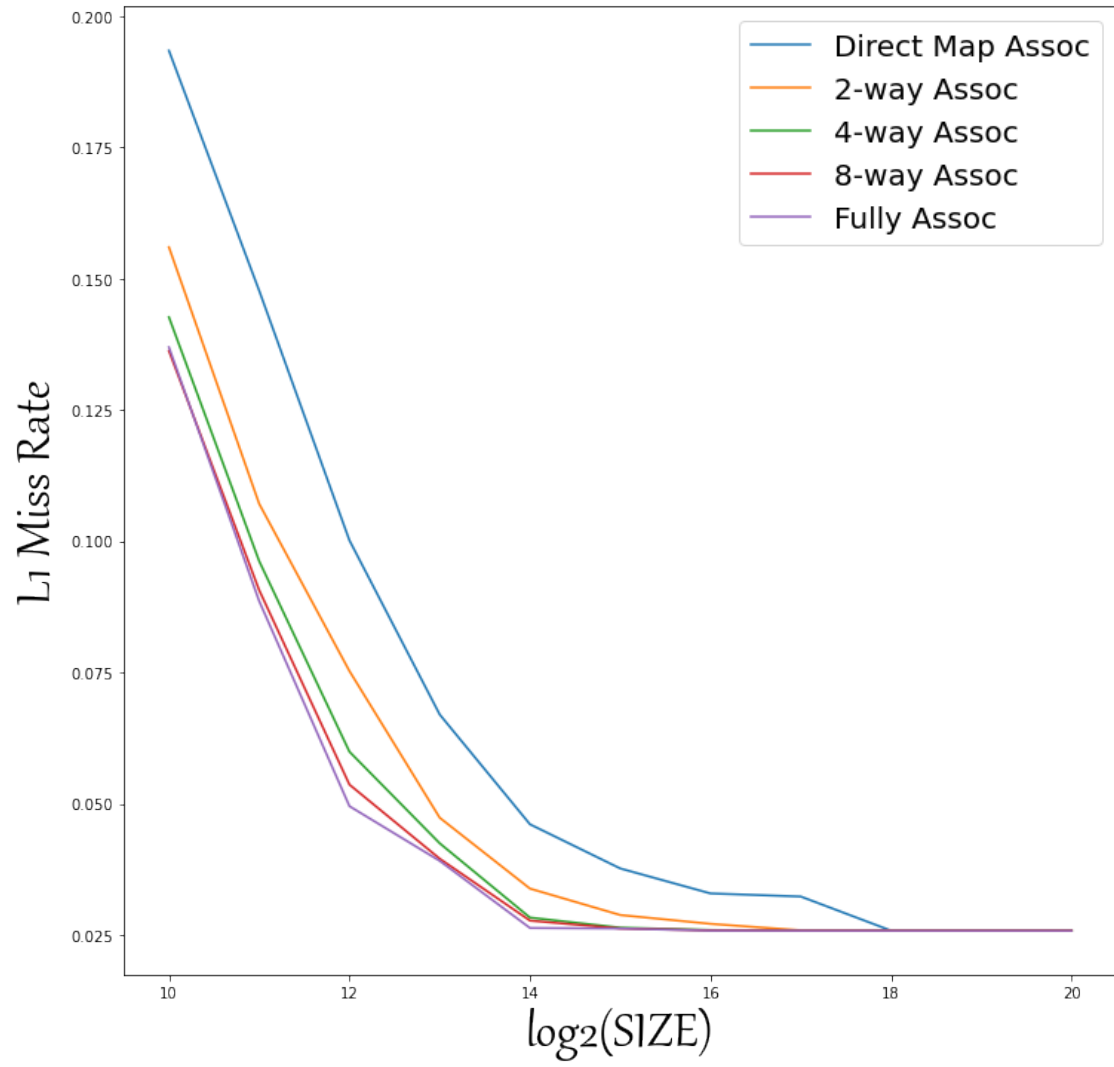
```

[4]: x = log1
y = [L1MissRates_Assoc1, L1MissRates_Assoc2, L1MissRates_Assoc4,↪
↪L1MissRates_Assoc8, L1MissRates_FullyAssoc ]
Label = ["Direct Map Assoc", "2-way Assoc", "4-way Assoc", "8-way Assoc",↪
↪"Fully Assoc"]

for i in range(len(y)):
    plt.rcParams['figure.figsize'] = [12, 12]
    #naming x axis
    plt.xlabel('log2(SIZE)', fontname="Gabriola", fontsize=40)
    # naming the y axis
    plt.ylabel('L1 Miss Rate', fontname="Gabriola", fontsize=40)
    plt.plot(x, y[i], label = Label[i], linestyle="-")

plt.legend(fontsize=20)
plt.show()

```



2.1 1. Discuss trends in the graph. For a given associativity, how does increasing cache size affect miss rate? For a given cache size, what is the effect of increasing associativity?

2.1.1 As it can be seen in the graph each line represents a specific associativity tracing it to different cache sizes. Given that the block size, replacement and inclusion policies are constant in all of them, this allows us to concentrate on effects of increased cache size with respect to each associativity. As it can be seen increasing cache size, decreases the miss rate so they have inverse relationship with one another. This fundamentally makes sense, the more cache we have the more data we can store in it. Hence, the less the likelihood of miss rates. However adding cache size at some point does not help the miss rates as much as it can be seen in the figure at some point miss rates do not change much by adding more cache and they turn into almost constant lines. which in turn is dependent on the number of Compulsory Misses

2.1.2 for a given cache size we would consider a vertical line through the graph. As it can be seen the higher the number of Associativity the lower the miss rates. So number of Associativity also has inverse correlation with cache size. That said the benefits of adding associativity at some point start having diminishing effects. It can be seen 2-way, 4-way, 8-way and full Associativity do almost merge at higher cache sizes and the 4-way, 8-way and fully Associate caches do follow each other very narrowly even at smaller cache sizes. So at some point adding more Associativity does not necessarily yield in much better miss rates

2.2 2. Estimate the compulsory miss rate from the graph.

2.2.1 All the graphs converge to 0.025820 so that would be the estimated compulsory miss rate. because that is the time the cache will minimally need to fetch the data originally needed from memory

2.3 3. For each associativity, estimate the conflict miss rate from the graph.

2.3.1  $\text{conflict miss} = \text{missrate} - \text{compulsory miss rate}$

```
[5]: cumPulsoryMissRate = 0.025820
    for i in range(len(y)):
        ConflictMissRate = y[i][0]-cumPulsoryMissRate
        print (Label[i], "Miss Rate = " , ConflictMissRate)
```

Direct Map Assoc Miss Rate = 0.16763999999999998

2-way Assoc Miss Rate = 0.13021

4-way Assoc Miss Rate = 0.11688

8-way Assoc Miss Rate = 0.11045

Fully Assoc Miss Rate = 0.11114

### 3 GRAPH #2 (no additional simulations with respect to GRAPH #1)

3.1 Same as GRAPH #1, but the y-axis should be AAT instead of L1 miss rate. Discussion to include in your report:

3.2 (Note: The value of Access time of 1KB 8-way Cache with block size 32 is not available in Cacti Table, you can ignore this point in your graph.)

3.3 1. For a memory hierarchy with only an L1 cache and BLOCKSIZE = 32, which configuration yields the best (i.e., lowest) AAT?

3.3.1  $AAT = 1 + 1 \text{ --- } ty$

```
[6]: y = [L1MissRates_Assoc1, L1MissRates_Assoc2, L1MissRates_Assoc4, L1MissRates_Assoc8, L1MissRates_FullyAssoc ]
MissPenalty = 100

## from excell
AccessTime1KB = [0.114797, 0.140329, 0.14682, 0.14682, 0.155484]
AccessTime2KB = [0.12909, 0.161691, 0.154496, 0.180686, 0.176515]
AccessTime4KB = [0.147005, 0.181131, 0.185685, 0.189065, 0.182948]
AccessTime8KB = [0.16383, 0.194195, 0.211173, 0.212911, 0.198581]
AccessTime16KB = [0.198417, 0.223917, 0.233936, 0.254354, 0.205608]
AccessTime32KB = [0.233353, 0.262446, 0.27125, 0.288511, 0.22474]
AccessTime64KB = [0.294627, 0.300727, 0.319481, 0.341213, 0.276281]
AccessTime128KB = [0.3668, 0.374603, 0.38028, 0.401236, 0.322486]
AccessTime256KB = [0.443812, 0.445929, 0.457685, 0.458925, 0.396009]
AccessTime512KB = [0.563451, 0.567744, 0.564418, 0.578177, 0.475728]
AccessTime1024KB = [0.69938, 0.706046, 0.699607, 0.705819, 0.588474]

AccessTimes = [AccessTime1KB ,
                AccessTime2KB ,
                AccessTime4KB ,
                AccessTime8KB ,
                AccessTime16KB,
                AccessTime32KB,
                AccessTime64KB,
                AccessTime128KB,
                AccessTime256KB,
                AccessTime512KB,
                AccessTime1024KB,
                ]

ATT = []

for i in range (5):
    A =[]
    for j in range(11):
```

```

        x = AccessTimes[j][i] + y[i][j]*MissPenalty
        A.append(x)
    ATT.append(A)

```

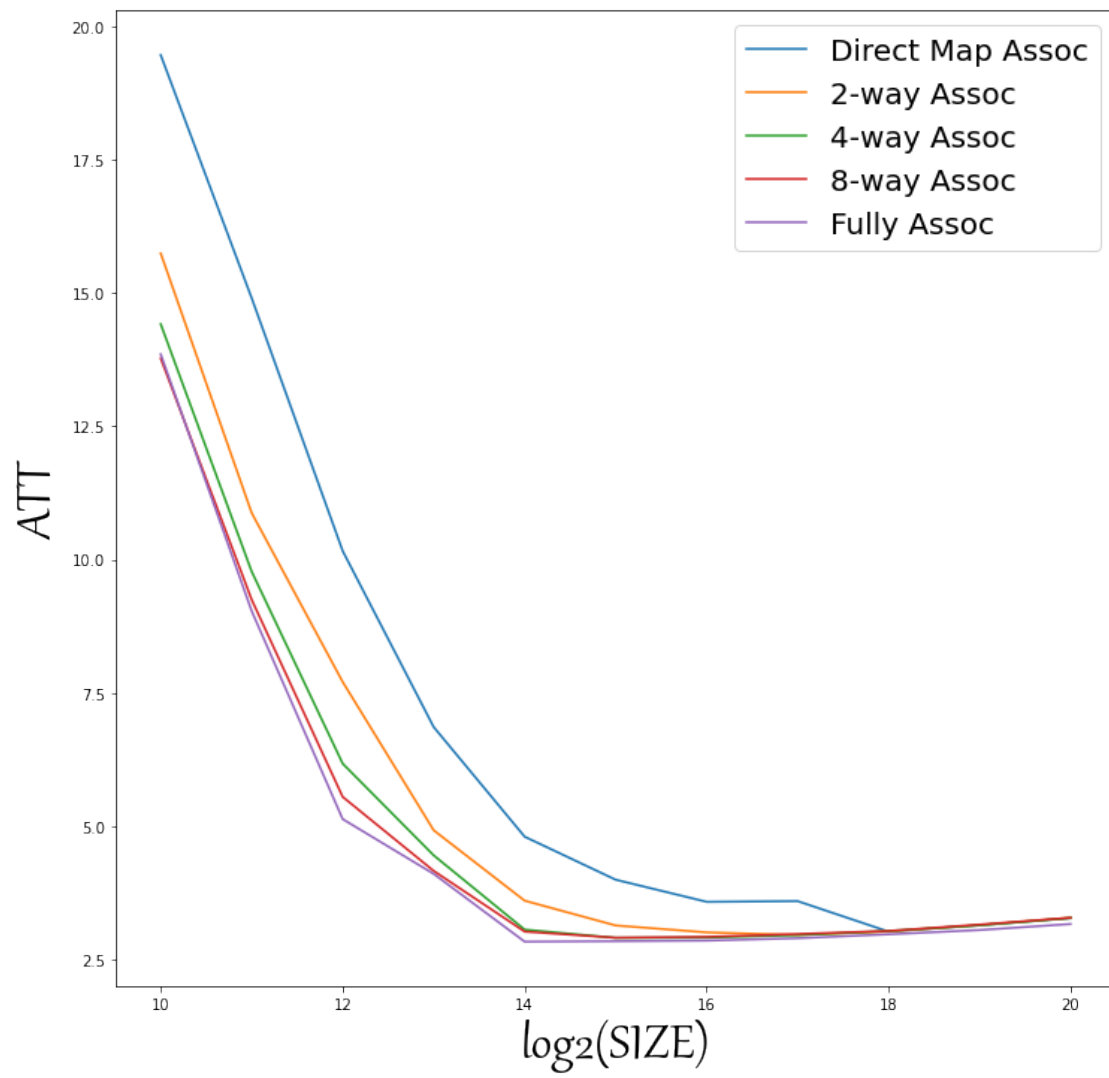
```

[7]: x = log1
     y = ATT
     Label = ["Direct Map Assoc", "2-way Assoc", "4-way Assoc", "8-way Assoc", "Fully Assoc"]

     for i in range(len(y)):
         plt.rcParams['figure.figsize'] = [12, 12]
         plt.xlabel('log2(SIZE)', fontname="Gabriola", fontsize=40)
         # naming the y axis
         plt.ylabel('ATT', fontname="Gabriola", fontsize=40)
         plt.plot(x, y[i], label = Label[i], linestyle="-")

     plt.legend(fontsize=20)
     plt.show()

```





3.3.2 Fully Associative yielding the best at cache size  $\log(\text{size}) = 14$  of 16384 bytes. As it can be seen the larger the size of the cache, the less ATT. At some point adding cache size starts having diminishing effects due to the fact that we are keeping unnecessary information in the cache. Sort of polluting the cache.

## 4 8.2. Replacement policy study

4.1 GRAPH #3 (total number of simulations: 27)

4.2 For this experiment:

4.3 o L1 cache: SIZE is varied, ASSOC = 4, BLOCKSIZE = 32.

4.4 o L2 cache: None.

4.5 o Replacement policy: varied

4.6 o Inclusion property: non-inclusive

4.7 Plot AAT on the y-axis versus  $\log_2(\text{SIZE})$  on the x-axis, for nine different cache sizes: SIZE = 1KB, 2KB, ... , 256KB, in powers-of-two. (That is,  $\log_2(\text{SIZE}) = 10, 11, \dots, 18$ .) The graph should contain three separate curves (i.e., lines connecting points), one for each of the following replacement policies: LRU, Pseudo-LRU, Optimal. All points for LRU replacement policy should be connected with a line, all points for Pseudo-LRU replacement policy should be connected with a line, etc.

4.8 Discussion to include in your report:

4.9 1. Discuss trends in the graph. Which replacement policy yields the best (i.e., lowest) AAT?

```
[8]: BlockSize = 32
size2 = []
Logs2=[]
for i in range(9):
    sizei = (2**i)*1024
    size2.append(sizei)
    log2 = math.log(sizei, 2)
    Logs2.append(log2)
    print("cashe ", i+1, "size of cashe in bytes: ", sizei, " and the log is",
    ↪log2 )

print(Logs2)
print(size2)
```

```
cashe 1 size of cashe in bytes: 1024 and the log is 10.0
cashe 2 size of cashe in bytes: 2048 and the log is 11.0
cashe 3 size of cashe in bytes: 4096 and the log is 12.0
cashe 4 size of cashe in bytes: 8192 and the log is 13.0
cashe 5 size of cashe in bytes: 16384 and the log is 14.0
cashe 6 size of cashe in bytes: 32768 and the log is 15.0
```

```

cashe 7 size of cashe in byes: 65536 and the log is 16.0
cashe 8 size of cashe in byes: 131072 and the log is 17.0
cashe 9 size of cashe in byes: 262144 and the log is 18.0
[10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0, 17.0, 18.0]
[1024, 2048, 4096, 8192, 16384, 32768, 65536, 131072, 262144]

```

```

[9]: L1MissRates_LRU = [0.142700, 0.096220, 0.059920, 0.042470, 0.028320, 0.026400, 0.025950, 0.025820, 0.025820]
L1MissRates_Pseudo_LRU = [0.143550, 0.096190, 0.060820, 0.042710, 0.028390, 0.026410, 0.025910, 0.025820, 0.025820]
L1MissRates_Optimal = [0.106090, 0.069360, 0.046220, 0.034650, 0.026850, 0.025980, 0.025820, 0.025820, 0.025820]

```

```

[10]: MissPenalty = 100
AccessTime1KB4 = [0.114797, 0.140329, 0.14682, 0.14682 ,0.155484]
AccessTime2KB4 = [0.12909, 0.161691, 0.154496, 0.180686, 0.176515]
AccessTime4KB4 = [0.147005, 0.181131, 0.185685, 0.189065, 0.182948]
AccessTime8KB4 = [0.16383, 0.194195, 0.211173, 0.212911, 0.198581]
AccessTime16KB4 = [0.198417,0.223917,0.233936, 0.254354, 0.205608]
AccessTime32KB4 = [0.233353, 0.262446, 0.27125, 0.288511, 0.22474]
AccessTime64KB4 = [0.294627, 0.300727, 0.319481, 0.341213, 0.276281]
AccessTime128KB4 = [0.3668, 0.374603, 0.38028, 0.401236, 0.322486]
AccessTime256KB4 = [0.443812, 0.445929, 0.457685, 0.458925, 0.396009]

AccessTimes2 = [AccessTime1KB ,
                AccessTime2KB ,
                AccessTime4KB ,
                AccessTime8KB ,
                AccessTime16KB,
                AccessTime32KB,
                AccessTime64KB,
                AccessTime128KB,
                AccessTime256KB,

                ]

y2 = [L1MissRates_LRU, L1MissRates_Pseudo_LRU,L1MissRates_Optimal]

ATT2 = []


for i in range (3):
    A2 =[]
    for j in range(9):
        x = AccessTimes2[j][2] + y2[i][j]*MissPenalty
        A2.append(x)
    ATT2.append(A2)

```

```
print (ATT2)
```

```
[[14.41682, 9.776496, 6.177685, 4.4581729999999995, 3.065936, 2.91125,  
2.9144810000000003, 2.96228, 3.039685], [14.50182, 9.773496, 6.267685,  
4.4821729999999995, 3.072936, 2.9122500000000002, 2.910481, 2.96228, 3.039685],  
[10.75582, 7.0904960000000001, 4.807685, 3.676173, 2.918936, 2.86925, 2.901481,  
2.96228, 3.039685]]
```

```
[11]: x = Logs2  
y = ATT2  
Label = ["LRU", "Pseudo_LRU", "Optimal"]  
  
for i in range(len(y)):  
    plt.rcParams['figure.figsize'] = [12, 12]  
    plt.xlabel('log2(SIZE)', fontname="Gabriola", fontsize=40)  
    # naming the y axis  
    plt.ylabel('ATT', fontname="Gabriola", fontsize=40)  
    plt.plot(x, y[i], label = Label[i], linestyle="-")  
  
plt.legend(fontsize=20)  
plt.show()
```



4.9.1 As it can be seen the larger the size of the cache, the less ATT . Among the policies, optimal policy gives us the best ATT. Which is understandable because in theory that is the best yield in terms of cache performance. Optimal Policy, Cache size of  $\log(\text{size}) = 14 = 16384$  Bytes. As it can be seen the larger the size of the cache, the less ATT. At some point adding cache size starts having diminishing effects due to the fact that we are keeping unnecessary information in the cache. Sort of polluting the cache.

#### 4.10 8.3. Inclusion property study

4.10.1 GRAPH #4 (total number of simulations: 12)

4.10.2 For this experiment:

4.10.3 o L1 cache: SIZE = 1KB, ASSOC = 4, BLOCKSIZE = 32.

4.10.4 o L2 cache: SIZE = 2KB – 64KB, ASSOC = 8, BLOCKSIZE = 32.

4.10.5 o Replacement policy: LRU

4.10.6 o Inclusion property: varied

4.10.7 Plot AAT on the y-axis versus  $\log_2(\text{L2 SIZE})$  on the x-axis, for six different L2 cache sizes: L2 SIZE = 2KB, 4KB, ... , 64KB, in powers-of-two. (That is,  $\log_2(\text{L2 SIZE}) = 11, 12, \dots, 16$ .) The graph should contain two separate curves (i.e., lines connecting points), one for each of the following inclusion properties: non-inclusive and inclusive. All points for non-inclusive cache should be connected with a line, all points for inclusive cache should be connected with a line.

4.10.8 Discussion to include in your report:

4.10.9 1. Discuss trends in the graph. Which inclusion property yields a better (i.e., lower) AAT?

```
[12]: BlockSize = 32
size3 = []
Logs3=[]
for i in range (6):
    sizei3 = (2**(i+1)*1024)
    size3.append(sizei3)
    log3 = math.log(sizei3, 2)
    Logs3.append(log3)
    print("cashe ", i+1, "size of cashe in bytes: ", sizei3, " and the log is",
    ↪log3 )

print(Logs3)
```

```
cashe 1 size of cashe in bytes: 2048 and the log is 11.0
cashe 2 size of cashe in bytes: 4096 and the log is 12.0
cashe 3 size of cashe in bytes: 8192 and the log is 13.0
cashe 4 size of cashe in bytes: 16384 and the log is 14.0
cashe 5 size of cashe in bytes: 32768 and the log is 15.0
```

cashe 6 size of cashe in bytes: 65536 and the log is 16.0  
 [11.0, 12.0, 13.0, 14.0, 15.0, 16.0]

```
[13]: L1MissRates_non_inclusive = [0.142700, 0.142700, 0.142700, 0.142700, 0.142700,
    ↪0.142700]
    L2MissRates_non_inclusive = [0.629713, 0.376174 , 0.277715, 0.194114, 0.183952,
    ↪0.181430]

    L1MissRates_inclusive = [0.145230, 0.143090, 0.142770, 0.142700, 0.142700, 0.
    ↪142700]
    L2MissRates_inclusive = [0.649108, 0.377944, 0.278000, 0.194114, 0.183952, 0.
    ↪181430]
```

$$5 \quad ATT = 1 + 1 \left( 2 + 2 \quad \_ \quad \right)$$

```
[18]: AccessTimes3 = [
    AccessTime2KB ,
    AccessTime4KB ,
    AccessTime8KB ,
    AccessTime16KB,
    AccessTime32KB,
    AccessTime64KB,
    ]

    AccessTime1KB

    y3 = [L2MissRates_non_inclusive, L2MissRates_inclusive]
    z3= [L1MissRates_non_inclusive, L1MissRates_inclusive ]

    ATT3 = []
    for i in range (2):
        A3 =[]
        for j in range(6):
            x = AccessTime1KB[2]+ (z3[i][j] * (AccessTimes3[j][3]+
    ↪(y3[i][j]*MissPenalty)))
            A3.append(x)
        ATT3.append(A3)


    print(ATT3)
```

```
[[9.1586084022, 5.5418025555, 4.140195449699999, 2.9531230957999997,
2.8129855596999995, 2.7845171951], [9.600056511779998, 5.58187400685,
4.146223303470001, 2.9531230957999997, 2.8129855596999995, 2.7845171951]]
```

```
[19]: x = Logs3
y = ATT3
Label = ["non_inclusive", "inclusive"]

for i in range(len(y)):
    plt.rcParams['figure.figsize'] = [12, 12]
    plt.xlabel('log2(SIZE)', fontname="Gabriola", fontsize=40)
    # naming the y axis
    plt.ylabel('ATT', fontname="Gabriola", fontsize=40)
    plt.plot(x, y[i], label = Label[i], linestyle="-")

plt.legend(fontsize=20)
# plt.legend()
plt.show()
```



```
[20]: ATT3[0] == ATT3[1]
```

```
[20]: False
```

- 5.1 lines seem very close but they are not the same as it can be seen above with comparing the np.arrays. Non inclusive has slightly better ATT. The larger the cache gets the ATT decreases. At some point again adding cache size just has diminishing effects as it can be seen ATT after  $\log(\text{size}) = 14$  starts flattening to almost a line and after  $\log(\text{size}) = 15$  its almost horizontal line.